

Content Disclaimer

This document contains examples of toxic, hateful, and offensive language generated for research and evaluation purposes only. The included text samples are public and are not endorsed by the author, organization, or reviewers. They are used solely to demonstrate adversarial testing, evaluate model robustness, and study mitigation strategies against harmful content. Viewer discretion is advised. If you may be sensitive to profanity, slurs, or hate speech, please proceed with caution.

Take-Home Assignment: Mini Red-Team Generator

Design and implement a model that could act as an adversarial red-team generator. The system should demonstrate how to **produce diverse toxic, harmful and hateful content**.

You can pick profanity, hate speech or harassment style generated contents. You are free to use any dataset as base or training set and any model which can achieve the goal.

We want to learn:

1. How do you approach this problem with a small dataset and a working model, why do you make those choices?
2. How do you evaluate your approach? Are the generated contents from your model good? Is this approach better than generic generative models?
3. The trade offs you made, and how you communicate them to the team?
4. End-to-end runnable solution / model

Deliverables:

1. Your writeup
2. Your PoC code zip
3. Your workable test cases

Example model generated outputs:

(You can generate potential hate / harmful toxic content in any scope, this is just a very restrictive scope example quoted from [HateXplain](#) dataset)

Offensive Related Output: "talk about retarded what a stupid gif you appear to be the poster boy for it"

Hatespeech Related Output: "my bf hit me but it ok he hits like a faggot real"